

Parallel Fully Dynamic Maintenance of 2-Connected Components

Chirayu Anant Haryan[†], G Ramakrishna[†], Kishore Kothapalli[§] and Dip Sankar Banerjee[‡]

[†]Indian Institute of Technology Tirupati, Tirupati, 517 506, India.

[§]International Institute of Information Technology Hyderabad, 500 032, Telengana, India.

[‡]Indian Institute of Technology Jodhpur, Rajasthan 342 037, India.

{cs18s503, rama}@iittp.ac.in, kkishore@iiit.ac.in, dipsankarb@iiitj.ac.in

Abstract—Finding the biconnected components of a graph has a large number of applications in many other graph problems including planarity testing, computing the centrality metrics, finding the (weighted) vertex cover, coloring, and the like. Recent years saw the design of efficient parallel algorithms for this problem across sequential and parallel computational models. However, current algorithms do not work in the setting where the underlying graph changes over time in a dynamic manner via the insertion or deletion of edges.

Dynamic algorithms that obtain the biconnected components of a graph upon insertion or deletion of a single edge are known from over two decades ago. Parallel algorithms for this problem are not heavily studied. In this paper, we design parallel algorithms that obtain the biconnected components of a graph subsequent to the insertion or deletion of a batch of edges. Our algorithms hence will be capable of exploiting the parallelism adduced due to a batch of updates.

We implement our algorithms on an AMD EPYC 7742 CPU having 128 cores. Our experiments on a collection of 10 real-world graphs from multiple classes indicate that our algorithms outperform parallel state-of-the-art static algorithms.

I. INTRODUCTION

Analysis of large graphs is a computationally expensive task. With the growing importance of online social networks, research targeted towards fast computations of vital graph metrics is gaining widespread importance [11, 20, 9]. One of the challenges is to update vital graph metrics of massive graphs as the graph changes over time. For instance, consider a social network such as Facebook, which is on average, adding about 500,000 new users every day (about 6 new profiles every second). Such new users will result in adding more edges to the network in terms of “friends”. Even at a low rate of one new edge per each new user, we can assume that the network gets 144 million new edges every day. Similarly, the actions of users in “unfriending” results in the network losing some edges every day. Similar settings apply to graphs from other domains such as transportation networks, biological networks, and collaboration networks. We, therefore, note that several classes of real-world networks face *churn*.

With current graph sizes of the order of several million to billions of nodes and edges, even trivial metrics such as connectivity can take a tremendous amount of time. Hence, repeated processing of the entire graph is not feasible where the temporal properties of the graph play an essential role. Therefore, we need to employ *dynamic algorithms* that treat

the graphs not as a static entity but as a dynamic entity and can process updates to the graph without warranting a re-computation of the graph analytic of interest.

Typical computations on such networks such as community detection, clustering, recommendation systems, and obtaining centrality metrics need, therefore, to take the churn in the network into account. Computations listed above, and also other computations such as finding the (weighted) vertex cover, coloring, and the maximal clique [7] use primitives such as finding the 2-connected components of the underlying graph [8, 14] to gain practical efficiency.

In this paper, we study parallel algorithms for maintaining the biconnected components of a graph in a dynamic setting. In particular, we envisage two situations in which we need to process a batch of updates to the underlying graph in parallel. We first look at an *incremental* method of inserting new edges in an existing graph. Next, we study the *fully* dynamic case where edges can be added to the current graph or deleted from the current graph. A batch of updates consists of edges added to the current graph or a set of edges deleted from the current graph. This separation into insert-only batches and delete-only batches allows us to simplify the problem and avoid issues related to concurrency.

Despite such separation into insert-only batches and delete-only batches, the problem is still challenging. As Galil and Italiano [3] point out, the fully dynamic maintenance of biconnectivity is much more challenging than the fully dynamic connectivity maintenance. The difference stems from the observation that in fully dynamic maintenance of biconnectivity, the number of biconnected components may change by $O(n)$ per every edge inserted or deleted.

There are a few prior works dealing with dynamic algorithms [2, 21, 5]. In the sequential computing model, an effective algorithmic technique called *sparsification* proposed by Eppstein et al. [2] shows how to design efficient dynamic algorithms for several graph problems. The sparsification approach suggests transforming the work associated with updating a sequence of updates on subgraphs of successively increasing size. The subgraphs themselves are called *sparse certificates*.

However, the sparsification approach has some limitations in the context of our problem. Firstly, the sparsification approach is designed to process a single update. The current scale




Fig. 1. Performance of Sparsification.

of the networks and the rate of churn do necessitate using parallelization in dynamic graph algorithms to gain efficiency. Secondly, we observe that the practical advantage of sparsification in a parallel setting is limited. As Figure 1 shows, using a parallel adoption of sparsification for maintaining the biconnected components includes multiple BFS traversals and hence is not competitive with respect to static algorithms.

Liang et al. [21] show a PRAM algorithm for maintaining the biconnected components of a graph in the dynamic setting. This algorithm that performs a single edge insertion/deletion in $O(\log^2 n)$ time using $O(n\alpha(m, n)/\log n)$ processors¹. The results of Liang et al. extend to related problems such as maintaining the 2-edge connectivity, 3-edge connectivity, and 3-connectivity of a graph. The work and time bounds mentioned above rely on the sparsification approach and, hence, are likely to be inefficient in practice.

We note from the survey of the existing works with respect to dynamic parallel maintenance of bi-connected components, that there still exists scope for the design of efficient and practical algorithms.

A. Our Contributions and Results

In this paper, we design novel parallel algorithms for maintaining the biconnected components of an undirected graph in a dynamic setting. We design a parallel algorithm (cf. Algorithm 1) that maintains the biconnected components of a graph under a batch of insertions. According to the nomenclature of dynamic graph algorithms, this algorithm is an incremental algorithm. We then extend this algorithm (cf. Algorithm 2, Algorithm 3) to fully dynamic algorithms for maintaining the biconnected components of a graph under a batch of insertions and deletions. These algorithms requires us to use novel data structures to record the impact of various edges on the biconnected components of the graph.

Our algorithms work by identifying the affected fundamental cycles due to the update and performing independent parallel operations on all the affected fundamental cycles to obtain the biconnected components of the new graph.

¹The quantity $\alpha(m, n)$ denotes the inverse Ackermann function.

TABLE I
THROUGHPUT ACHIEVED BY OUR ALGORITHMS.

Algorithm/Operation	Insert	Delete	
		Non-Tree	Tree
Incremental	12 M	-	-
Fully Dynamic	63 K	42 K	1 K

Our study is motivated by the following factors: the practical possibility that in real-world networks with high churn, one can prepare insert-only and delete-only batches, the possibility of exploiting parallelism that one can adduce via a batch of updates, and the increasing size of the current generation networks (order of 1 billion nodes) and the availability of massive parallelism through modern generation hardware.

Experiments on a set of real-world (from [10] and [1]) graphs suggest that our algorithms Incremental Dynamic and Fully Dynamic offer a significant advantage compared to static algorithms. The throughput that our algorithms achieve on a multi-core CPU is summarized in Table I.

B. Related Work

Sequential algorithms for maintaining the biconnected components of a graph under single edge insertion or deletion are studied in several works, viz. Galil and Italiano [3], Henzinger and King [6]. The sparsification approach of Eppstein et al. [2] results in algorithms that can perform one update in $O(n)$ time. In the sparsification-based approach, Eppstein et al. [2] propose the idea that for several graph problems, one can perform the update on a sparse certificate of the original graph. The sparse graph has only $O(n)$ edges of the original graph. This approach has the advantage that the actual update is handled on a much smaller graph than the original graph.

Ramalingam [13] argues that the time taken to process an update be measured in terms of the impact the new edge has on the analytic. Ramalingam refers to this as the incremental complexity and shows that for many graph problems, one can design algorithms that work in time proportional to the incremental complexity of the problem.

Liang et al. [21] show an PRAM algorithm for maintaining 2-edge connectivity that can process a single update in $O(\log n \log(m/n))$ time using $O(n^{3/4})$ processors. Further, they also show that 2-vertex connectivity can be maintained in time $O(\log^2 n)$ using $O(n\alpha(2n; n) \cdot \log n)$ processors for a single update where $\alpha(\cdot)$ is the inverse Ackermann's function. They also show how to extend their results to the case of maintaining 3-edge and 3-vertex connectivity for a single update.

The sparsification approach [2] has been used in several recent parallel dynamic graph algorithms. Bhowmick and Das [19] in their recent work on maintaining dynamic connectivity. Khanda et al. [9] study the dynamic maintenance of single-source shortest paths in an evolving graph. The main idea from Khanda et al. [9] is to locate the affected vertex (vertices) and update the required shortest paths in a load-balanced manner. The work in [9] presents the batch parallel case for their single update case shown in [20].

McColl et al. [12] show how to handle graph connectivity in a dynamic setting on real-world graphs. They use GPUs for their experimental evaluation and show that all but a small minority of queries require very little processing time in updating the connectivity as the graph changes. Simsiri et al. [16] also study the connectivity problem on dynamic graphs but introduce a model for handling updates in bulk. On multicore CPUs, they show how to process a stream of edges to the added to the graph in small batches, and the effect of each batch can be processed in logarithmic time in parallel and uses work that is nearly linear in the size of the batch. Their work, however, addresses only the incremental version of the problem.

Jamour et al. [8] use the idea of Green et al. [4] that identifies *affected* vertices due to a single edge insertion or deletion and run BFS from the affected vertices to update the betweenness-centrality values of nodes in a graph. This approach is extended by Regunta et al. [15] to handle the case of a batch update and also to the case of updating the closeness-centrality values of nodes in a graph.

II. CUT VERTICES VIA UNSAFE COMPONENTS

We start this section with the standard graph-theoretic terminology. Later, we introduce the notion of safe and unsafe components in LCA-graphs. Finally, we present the necessary and sufficient conditions for a vertex to be a cut vertex using the notion of unsafe components.

We use G to denote an unweighted and undirected graph. A vertex v in G is a *cut vertex* if $G - v$ is disconnected. Similarly, an edge e in G is a *cut edge* if $G - e$ is disconnected. A graph with no cut vertex is called *biconnected*. A maximal biconnected subgraph of G is referred to as a *biconnected component* of G . The number of neighbours of a vertex v in G is called the *degree* of v , denoted by $d(v)$. Let T be a rooted spanning tree of G , and r be the root of T . An edge in G is a *non-tree edge* if does not appear in T . A non-tree edge $e = (u, v)$ along with the tree path between u and v forms a cycle, C_e known as a *fundamental cycle*. For a non-tree edge $e = (u, v)$ of G , we use $C_{(u,v)}$ or C_e to denote the fundamental cycle formed with the tree path between u and v and the non-tree edge (u, v) . A vertex in T is an LCA vertex if it is the least common ancestor for the end points of some non-tree edge e ; for simplicity, we say that v is LCA for edge e as well as fundamental cycle C_e . We define the notion of base vertices and base edges for a fundamental cycle C_e whose LCA is w as follows: The set $B = \{b \mid b \text{ is a child of } w \text{ in } T \text{ and } b \in V(C_e)\}$ is the set of *base vertices*. Note that $|B| \leq 2$. In case $|B| = 2$, the edge between the two vertices in B is the *base edge* of C_e . For a vertex w in T , the set of base vertices and base edges of all the fundamental cycles whose LCA is w forms a graph G_w , referred to as an LCA-graph. In case w is not an LCA, we can observe that G_w is an empty graph. A vertex u in T is a *safe vertex*, if u is part of a fundamental cycle whose LCA vertex is not equal to the parent of u ; otherwise u is *unsafe*. For any two vertices x and y in an LCA-graph G_w , we observe that there exists a path between

x and y in G_w if and only if there exists path between x and y in G , without going through w . This observation inspires to define the notion of safe and unsafe components in G_w . A component H in G_w is *safe* if it has at least one safe vertex; otherwise H is *unsafe*. The importance of presence of a safe vertex in a component of an LCA-graph is described in the following lemma.

Lemma 2.1: Let T be a rooted spanning tree of G , and w be an LCA vertex in T . Let x be a vertex in G_w and y be the parent of w . There is a path between x and y in $G - w$ if and only if x is reachable to a safe vertex in G_w without going through w .

The notion of unsafe components in LCA-graphs helps to derive the necessary and sufficient condition for a vertex to be a cut vertex.

Theorem 2.2: Let T be a rooted spanning tree of an unweighted graph G , r be the root of T , and w be an arbitrary vertex in T . w is a cut vertex in G if and only if one of the following holds true. a) $w \neq r$ and G_w has an unsafe component, b) $w = r$ and G_w has at least two components, and c) w is incident to a cut edge and $d(w) \geq 2$.

Proof: We first prove the sufficient condition for a vertex to be cut vertex. We first consider the case in which $w \neq r$ and G_w has an unsafe component H . Let x be an arbitrary vertex in H , which is unsafe. The vertex x can not reach any safe vertex without going through w , because all the vertices in H are unsafe. Therefore, there is no path between x and the parent of w in $G - w$ due to Lemma 2.1, and thus w is a cut vertex. We now consider the second case in which $w = r$, and G_w has at least two components. Let H_1 and H_2 be two components in G_w . Due to the construction of G_w , there are no non-tree edges crossing H_1 and H_2 . Hence, there is no path in $G - w$ between a vertex in H_1 and a vertex in H_2 , which implies that w is a cut vertex. In the last case, w is incident to a cut edge (w, x) , and let y be a neighbour of w such that $x \neq y$. The removal of w separates x from y , and thus w is a cut vertex.

We now prove the contrapositive of the necessary condition. In other words, we assume that the three conditions a), b), and c) are false, and show that w is not a cut vertex in G . If $d(w)$ is 1, then w is not a cut-vertex. Accordingly, in the rest of the proof, we assume that $d(w) \geq 2$ and w is not incident to a cut edge, because the condition c) is false. We first consider the case that $w \neq r$. Let H be a safe component in G_w and x be a safe vertex in H . Every vertex in H can reach x without going through w , as H is connected. Since every component in G_w is safe, from Lemma 2.1 the children of w who appear in G_w can reach the parent of w in $G - w$. Let y be a child of w such that y is not in G_w . Since w is not incident to a cut edge and y is not in G_w , (w, y) belong to a fundamental cycle whose LCA is not equal to w . Therefore y is a safe vertex. Thus, the children of w in T do not disconnect from the parent of w in $G - w$. The neighbours of w that are not children of w are any way reachable from the parent of w in $G - w$. Thus, w is not a cut vertex. We now consider the case that $w = r$. As the number of components in G_w is one, for every two vertices

Algorithm 1 INCREMENTAL-DYNAMIC

Input: The incremental dynamic data-structure \mathcal{D} of G , a batch \mathcal{B} of edges to be inserted in G .

Task: Update \mathcal{D} and find the cut vertices and cut edges in $G + \mathcal{B}$.

```
1: for each edge  $e = (u, v)$  in  $\mathcal{B}$  do in parallel
2:    $(V_*, E_*) = \text{RETRIEVE-BASE-VERTICES-EDGES}(T, e)$ 
3:    $F_* = \text{UPDATECONNECTCOMP}(F_*, V_*, E_*)$ 
4:   for each vertex  $v$  in  $F_*$  do in parallel
5:      $\text{safe}[\text{rep}[v]] = \text{safe}[\text{rep}[v]] \text{ or } \text{safe}[v]$ 
6:   for each vertex  $v$  in  $G$ , s.t  $\text{rep}[v] = v$  do in parallel
7:      $\text{hasUcomp}[p[v]] = \text{hasUcomp}[p[v]] \text{ or } \neg \text{safe}[v]$ 
8:      $\text{numCcomp}[p[v]] = \text{numCcomp}[p[v]] + 1$ 
9:   for each vertex  $w$  in  $G$  do in parallel
10:     $\text{cutVertex}[w] = (w \neq r \text{ and } \text{hasUcomp}[w] \text{ is true})$ 
    or  $(w = r \text{ and } \text{numCcomp}[r] > 1)$  or  $(w \text{ is incident to}$ 
     $\text{a cut-edge and } d(w) > 1)$ 
```

in G_w , there is a path in G without going through w . Thus, the removal of w does not disconnect the graph, and hence w is not a cut vertex. ■

III. A PARALLEL INCREMENTAL DYNAMIC ALGORITHM

In this section, we describe a parallel algorithm for inserting a batch \mathcal{B} of edges, to a current graph G and update the underlying data-structure to retrieve the cut vertices and cut edges in $G + \mathcal{B}$. Algorithm 1 shows a high level view of our incremental dynamic algorithm. Interestingly, it turns out that we can also obtain an algorithm for obtaining the biconnected components of the updated graph using Algorithm 1. This is explained in Section III-A.

Key Idea: Based on Theorem 2.2, we mark a vertex v in a graph G as a cut-vertex if the following holds. I) v is incident to a cut-edge (u, v) such that $d(v) \geq 2$. II) v is an LCA for the end vertices of a non-tree edge and G_v contains at least one unsafe component. The key idea in our dynamic algorithms is to update whether an edge is a cut-edge or not and update the components along with their safeness value in LCA graphs associated with LCA vertices, by traversing through the affected fundamental cycles.

For any two vertices x and y in G , the corresponding LCA-graphs G_x and G_y are vertex disjoint. This observation triggers us to introduce the notion of *global* LCA-graph G_* , where G_* is the disjoint union of G_x for every vertex x in G . The global LCA-graph allows to perform edge level parallelism on all edges in G_* .

We now present various components of the incremental dynamic data-structure $\mathcal{D} = (r, p, \text{safe}, \text{cutVertex}, \text{cutEdge}, \text{rep}, \text{hasUcomp}, \text{numCcomp})$. A spanning tree T of G rooted at r is stored using an integer array $p[\]$, where $p[v]$ stores the parent of v in T . $\text{safe}[\]$ is a boolean array on the vertices, and we set $\text{safe}[v]$ to true if and only if v is a safe vertex. Similarly, $\text{cutVertex}[\]$ and $\text{cutEdge}[\]$ are the boolean arrays on vertices and edges to indicate whether a given vertex or an edge is a cut vertex or a cut-edge, respectively. We

ensure that all the vertices in a component of G_* have the same representative. To this end, we use an integer array $\text{rep}[\]$ and $\text{rep}[v]$ stores the representative of v in F_* , where F_* is a set of rooted trees, and each tree corresponds to a component in G_* . $\text{hasUcomp}[w]$ is set to true if there G_w has an unsafe component; otherwise it is set to false. The number of connected components in G_w is stored in $\text{numCcomp}[w]$. **Description of the INCREMENTAL-DYNAMIC Algorithm.** Our incremental dynamic algorithm allows us to insert a batch of edges and compute the cut vertices and cut edges in the updated graph. Algorithm 1 details the major steps of our approach.

When a batch \mathcal{B} of edges are inserted, we traverse each fundamental cycle formed with an edge e in \mathcal{B} in parallel, using $\text{RETRIEVE-BASE-VERTICES-EDGES}(T, e)$. This function identifies the set V_* of base vertices and E_* of base edges of C_e , which are to be added to G_* . Also, the applicable vertices are marked as safe while traversing a fundamental cycle. Then, all the connected components F_* are updated by inserting V_* and E_* to G_* using $\text{UPDATECONNECTCOMP}(F_*, V_*, E_*)$. Later, all the vertices in F_* propagate their safeness to their representatives. A representative or a component is marked as safe if it has at least one safe vertex in Line 5. Further, each representative x in F_w (F_*) propagates whether it is a safe component or not to w (parent of x in T). The existence of an unsafe component in G_w is marked in $\text{hasUcomp}[w]$ in Line 7. The number of connected components in G_w is maintained in $\text{numCcomp}[w]$ in Line 8. Finally, the three sufficient conditions from Theorem 2.2 are applied in Line 10, to mark the cut vertices in the updated graph.

RETRIEVE-BASE-VERTICES-EDGES(T, e): We first identify the LCA vertex w of e . Later, we go through all the tree edges $e = (x, y)$ in the fundamental cycle C_e , where x is the parent of y and execute the following: $\text{cutEdge}[e]$ is updated to *false* as e belongs to a cycle b) If $x \neq w$, $\text{safe}[y]$ is updated to *true*. Further, the base vertices and base edges of C_e are added to V_* and E_* , respectively.

UPDATECONNECTCOMP(F_*, V_*, E_*): The purpose of this function is to update the connected components F_* of G_* by including V_* and E_* . We achieve this by using the parallel algorithm to identify connected components [18]. This algorithm maintains the set-disjoint-union data structure in parallel. The two key operations involved to find the connected components of G_* are namely *grafting* and *short-cutting*. The set of vertices of a connected component in F_* are represented using a rooted tree. For each vertex v in G , we perform the following in parallel to include new bases vertices in F_* as singleton sets: if $v \in V_*$ and $v \notin F_*$, $\text{rep}[v] = v$. When an edge in E_* is added across two components, the grafting operation joins the corresponding rooted trees. The other operation short-cutting converts every rooted tree to a star graph. The edges in E_* across the trees are treated as active edges for the next iteration. The above two operations grafting and short-cutting, happens on active edges in every iteration until there are no more active edges. The grafting happens on all active edges in parallel, whereas the short-

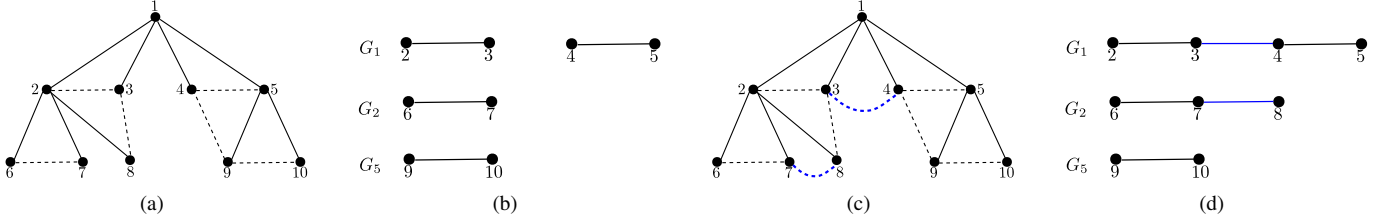


Fig. 2. (a) A graph G and a rooted spanning tree T of G are shown. Solid edges are edges of T , and dashed edges are non-tree edges of T . (b) Non-empty LCA-graphs are shown. (c) The edges (3, 4) and (7, 8) are inserted. (d) LCA-graphs are updated.

cutting works on each vertex in F_* in parallel.

Now, we illustrate Algorithm 1 using Fig 2. The vertices 8 and 9 are safe due to the fundamental cycles $C_{(3,8)}$ and $C_{(4,9)}$, respectively. As G_1 has two components $\{2, 3\}$ and $\{4, 5\}$, and 1 is a root vertex, 1 is a cut vertex due to condition-(b) in Theorem 2.2. By applying condition-(a) in Theorem 2.2, we can declare that 2 is a cut vertex, because G_2 has an unsafe component $\{6, 9\}$ and 2 is not a root vertex. G_5 has only one safe component $\{9, 10\}$, and no cut edge is incident to 5, and hence 5 is not a cut vertex. When (3, 4) is inserted, the fundamental cycle $C_{(3,4)}$ is traversed and edge (3, 4) is added to G_1 . Now G_1 has only one component $\{1, 2, 3, 4\}$ and hence 1 is not a cut vertex. When (7, 8) is inserted, $C_{(7,8)}$ is traversed and (7, 8) is added to G_2 . Since G_2 has only one safe component, we can declare that 2 is not a cut vertex.

A. Parallel Biconnected Components Algorithm

An edge e in T is a cut edge if $cutEdge[e]$ is true, whereas all the non-tree edges are known to be non-bridges. All the cut edges are trivial biconnected components, which can be computed using $cutEdge[\]$. We observe that each non-trivial biconnected component in G is corresponding to an unsafe component in G_* . For each representative u of an unsafe component in G_* , we perform restricted variant of the BFS algorithm starting at u in parallel. In the restricted BFS algorithm, we avoid inserting cut vertices into the queue. Each restricted BFS precisely traverses through the edges of a biconnected component of G .

B. A Static Algorithm

We now derive a static algorithm to find the cut vertices, cut edges and the biconnected components in a graph, using the incremental dynamic algorithm. Given an unweighted graph G , we first obtain a spanning tree T of G , and initialize various components in our data-structure \mathcal{D} as follows. Every vertex in T except the leaves are marked as cut vertices. All the edges in T are marked as cut edges. Every vertex in T is marked as unsafe. We then run Algorithm 1 on a batch \mathcal{B} formed with the edges in $G - T$.

IV. A FULLY DYNAMIC ALGORITHM

In this section, we first introduce the notion of partially and completely affected LCA vertices, which helps to avoid redundant work. Later, we describe individual components of

our dynamic data-structure. Finally, we present fully dynamic algorithms to support inserting/deleting a batch of edges.

A vertex w is called as *completely affected* if the topology of G_w is changed due to insertion/deletion of edges to and from G . Similarly, a vertex w is termed as *partially affected* if the $safeness(v)$ becomes zero from non-zero or vice versa. The completely and partially affected vertices help to avoid recomputing the connected components and recounting the number of unsafe components in G_w .

We now present various components of the fully dynamic data-structure $\mathcal{D} = (r, p, safe, cutVertex, cutEdge, rep, hasUcomp, numCcomp, Safeness, sup, pList, cList, LCA-strength, weight)$ that are not part of the incremental dynamic data-structure. The number of fundamental cycles that support a vertex v as safe is called *safeness* of v and is stored in $safeness[v]$. For a tree edge e , $sup[e]$ is a set of fundamental cycles containing e . $pList$ and $cList$ are sequences of partially and completely affected LCA vertices. $LCA-strength[v]$ stores the number of non-tree edges whose LCA is v .

Description of FULLY-DYNAMIC-INSERT-BATCH Algorithm. The pseudo-code of our fully dynamic parallel algorithm, which allows us to insert a batch of edges and maintain the cut vertices and cut edges is shown in Algorithm 2. When a batch \mathcal{B} of edges are inserted, we traverse each fundamental cycle formed with an edge in \mathcal{B} in parallel, using $INSERT-FUND-CYCLE(e, pList, cList)$. This function updates the topology of the necessary LCA graph, updates the safeness of applicable vertices in C_e , and finally all the partially affected and completely affected vertices are appended to $pList$ and $cList$, respectively. All the duplicates from $pList$ and $cList$ are removed in the next step. We then recompute the connected components in G_w using $CONNECTCOMP(G_w)$ for each w in $cList$ in parallel. Similarly, we recount the number of unsafe-components in G_w using $NUMUCOMP(G_w)$ for each vertex w in $pList \cup cList$ in parallel. Finally, we apply Theorem 2.2 to find the cut vertices in the updated graph.

INSERT-FUND-CYCLE($e, pList, cList$): The purpose of this function is to insert the non-tree edge e , traverse the fundamental cycle C_e , which is formed with T and e , and append the partially and completely affected LCA vertices to $pList$ and $cList$, respectively. We first identify the LCA w for the end points of e in T and increment $LCA-strength[w]$ by one. For each tree edge e' in C_e , we add e to $sup[e']$. For each vertex v in C_e , such that $p[v] \neq w$, we perform the following: increment $safeness[v]$ by one; If $p[v]$ is an

Algorithm 2 FULLY-DYNAMIC-INSERT-BATCH

Input: The fully dynamic data-structure \mathcal{D} of an unweighted graph G , a batch \mathcal{B} of edges to be inserted to G .

Task: Update \mathcal{D} and find the cut vertices and cut edges in $G + \mathcal{B}$.

Phase-I

- 1: **for** each edge (u, v) in \mathcal{B} **do in parallel**
- 2: INSERT-FUND-CYCLE($e, pList, cList$)

Phase-II

- 3: Remove duplicates in $pList$ and $cList$ in parallel
 - 4: **for** each vertex w in $cList$ **do in parallel**
 - 5: $F_w = \text{CONNECTCOMP}(G_w)$
 - 6: **for** each vertex w in $cList$ and $pList$ **do in parallel**
 - 7: $hasUcomp[w] = \text{HASUCOMP}(G_w)$
 - 8: **for** each vertex w in G **do in parallel**
 - 9: $cutVertex[w] = (w \neq r \text{ and } hasUcomp[w] \text{ is true})$
 or ($w = r$ and $\text{NUMCONNECTCOMP}(G_r) > 1$) **or** (w is
 incident to a cut-edge and $d(w) > 1$)
-

LCA, and $safeness[v]$ is updated from zero to non-zero, then append $p[v]$ to $pList$. Later, we identify the base vertices and base edges of C_e . For each base vertex x of C_e , if x is not in G_w , then we add x to G_w and initialize $weight(x)$ with one; otherwise increment $weight(x)$ by one. Similarly, if the base edge e' of C_e is not in G_w , then e' is added to G_w and set $weight(e') = 1$; otherwise $weight(e')$ is incremented by one. Finally, if any base vertex or base edge is added to G_w , then w is append to $cList$ to note down that G_w is structurally updated.

DELETE-FUND-CYCLE($e, pList, cList$): The aim of this function is to delete the non-tree edge e , traverse the fundamental cycle C_e , and append the partially and completely affected LCA vertices to $pList$ and $cList$, respectively. The steps of this function are similar to that of INSERT-FUND-CYCLE. We decrement LCA-strength and $safeness$ by one rather than incrementing by one. Similarly, we delete edges from sup rather than adding. Also, decrement the weights of base vertices and base edges if their weight is at least 2; otherwise, delete them accordingly. Finally, $p[v]$ is added to $pList$ if $safeness[v]$ becomes zero. Likewise, w is added to $cList$ if G_w is structurally updated.

CONNECTCOMP(G_w): We run the BFS algorithm on G_w and finds a forest F_w of rooted trees, where each tree corresponds to a spanning tree of a component in G_w . During the BFS traversal in G_w starting at a vertex, say x , every newly visited vertex updates its parent to x . This way, all the vertices in a component are stored in the rooted star-graph representation and root vertex is called as a representative..

HASUCOMP(G_w): Every safe vertex x in G_w propagates true to the representative of x . In case the x does not receive true from the children and x is not safe, such a component is treated as unsafe, and x propagates the existence of an unsafe component to w .

NUMCONNECTCOMP(G_w): This function finds the number of connected components in G_w by running the BFS algorithm.

Description of FULLY-DYNAMIC-DELETE-BATCH Algorithm. Algorithm 3 supports to delete a batch \mathcal{B} of edges from the current graph and updates the cut vertices and cut edges dynamically. The batch \mathcal{B} can have both tree and non-tree edges. This algorithm primarily consists of three phases. Phase-I is responsible to delete non-tree edges and delete the corresponding base vertices and base edges from the affected LCA-graphs, using DELETE-FUND-CYCLE($e, pList, cList$). Similarly, Phase-II deletes tree edges in \mathcal{B} from T , connect the disconnected trees of T with non-tree edges, and updates the base vertices and base edges accordingly from the corresponding LCA-graphs. Phase-III computes the connected components in LCA-graphs, identifies the existence of unsafe components, and finally updates the cut vertices.

We now elaborate the steps of Phase-II. Every vertex in a fundamental cycle C_e except the base vertices of C_e and the LCA of C_e is a safe vertex. Accordingly, we can observe that the addition or deletion of a fundamental cycle affects the safe and unsafe components in LCA-graphs. If we delete a tree edge e' , LCA-graph of $G_{lca(e')}$ is affected, where $e \in sup(e')$. To track all these changes for all the tree edges to be deleted, we collect the $sup(e)$ for each tree edge e in \mathcal{B} in $globalSup$ in Line 3. Later all the fundamental cycles formed with the non-tree edges in $globalSup$ are deleted using DELETE-FUND-CYCLE($e, pList, cList$) in Line 5. Now, we delete each tree edge (x, y) in \mathcal{B} from T in parallel, by updating $p[y]$ to y , where x was the old parent of y . A few non-tree edges in $globalSup$ become tree edges to reconnect the components in T using RECONNECT($T, globalSup$) (Line 7) and each such non-tree edge is called *co-edge*. For all the rest of the non-tree edges in $globalSup$, we add new fundamental cycles and update the necessary data structures using INSERT-FUND-CYCLE($e, pList, cList$) in Line 9.

RECONNECT($T, globalSup$): We first construct a simple super graph \mathbb{G} in parallel, where $V(\mathbb{G}) = \{\text{root of a tree } t \mid t \text{ is a component in } T\}$ and $E(\mathbb{G}) = globalSup$. For every edge (t_1, t_2) in \mathbb{G} , we associate a *co-edge* (x, y) in $globalSup$ such that x and y are contained in trees rooted at t_1 and t_2 , respectively. Then a spanning tree \mathbb{T} of \mathbb{G} is obtained using the parallel BFS algorithm. Further, for each edge (u, v) in \mathbb{T} in parallel, we perform the following to reconnect the components in T : Obtain the co-edge (x, y) of (u, v) , where x is the parent of y . Reverse the path between z and y and update the necessary LCA-graphs using REVERSEPATH(z, y), where z is the root of the tree containing y . Finally update the parent of y to x to insert the co-edge (x, y) .

REVERSEPATH(z, y): This function aims to reverse the path between z and y and do the necessary updates on certain LCA-graphs due to affected non-tree edges. A non-tree edge e' that belongs to $sup(e)$ is called *affected* if e appears in the reversing path. We first reverse the path between z and y by updating the parents of the vertices involved in the path. We then walk from both the endpoints of an affected non-tree edge towards the root until we encounter a special vertex. The two special vertices found share an ancestor-descendant relationship. The ancestor vertex is the old LCA,

Algorithm 3 FULLY-DYNAMIC-DELETE-BATCH

Input: The fully dynamic data-structure \mathcal{D} of an unweighted graph G , a batch \mathcal{B} of edges to be deleted from G .

Task: Update \mathcal{D} and find the cut vertices and cut edges in $G - \mathcal{B}$.

Phase-I

- 1: **for** each non-tree edge e in \mathcal{B} **do in parallel**
- 2: DELETE-FUND-CYCLE($e, pList, cList$)

Phase-II

- 3: $globalSup = \cup_{e \in \text{Tree Edges in } \mathcal{B}} sup(e)$ in parallel
- 4: **for** each non-tree edge e in $globalSup$ **do in parallel**
- 5: DELETE-FUND-CYCLE($e, pList, cList$)
- 6: Remove all the tree edges in \mathcal{B} from T in parallel
- 7: RECONNECT($T, globalSup$)
- 8: **for** each non-tree edge e in $globalSup$ such that e is not a co-edge **do in parallel**
- 9: INSERT-FUND-CYCLE($e, pList, cList$).

Phase-III

- 10: Call **Phase-II** in Algorithm 2
-

and the descendant is the new LCA vertex. The neighbors of the identified LCA vertices in the walk and the special path are the corresponding base vertices. We then update LCA-graphs and LCA-strengths of old and new LCA vertices, along with the safeness values of old and new base vertices.

V. IMPLEMENTATION DETAILS AND OPTIMIZATIONS

In this section, we discuss the technique to obtain LCA in our proposed algorithm. Further, we explain the details about the maintenance of the LCA graphs. Finally, we provide the details to update the data structures for a race-free implementation.

A. Finding LCA

A key operation in both of our proposed algorithms is to find the LCA vertex of a given non-tree edge. In Algorithm 1, we use the level numbers obtained from T to find the LCA vertex. As the spanning tree changes in the Algorithm 3, we do not use the level numbers. Instead, we perform a walk towards the root alternatively from the endpoints of the non-tree edge. We mark the visited ancestors during the walk, and the first ancestor visited by both walks is the LCA vertex. We maintain an integer visited array and mark the visited vertex with a unique non-tree edge number. This allows us to reuse the same visited array for all the non-tree edges, without initializing multiple times. However, each thread maintains a local visited array to avoid race conditions. The former method is more efficient as it only traverses through the tree edges that are part of the fundamental cycle.

B. LCA Graphs

In Algorithm 1, we maintain the global LCA graph as an edge list. Each thread maintains a local edge-list to avoid serial append to a global list. For UPDATECONNECTCOMP, each thread processes its own list, thus avoiding the need to merge the lists. On the contrary, we maintain LCA-graphs explicitly

TABLE II
DATASETS OBTAINED FROM [10] AND [1]

Graph name	V	E
Road Networks		
roadNet-CA	1,965,206	2,766,607
roadNet-PA	1,088,092	1,541,898
europe_osm	50,912,018	108,109,320
Web Graph		
web-Stanford	281,903	2,312,497
web-Google	875,713	5,105,039
Product co-purchasing Networks		
amazon0302	262,111	1,234,877
Citation networks		
cit-Patents	3,774,768	16,518,948
Networks with ground-truth communities		
com-Orkut	3,072,441	117,185,083
Social networks		
soc-LiveJournal1	4,847,571	68,993,773
LAW		
webbase-2001	118,142,155	1,019,903,190

in Algorithms 2 and 3. The edges incident to a vertex in a LCA-graph are stored in a *Set* container from STL. This is efficient to allow frequent search, insert, and delete operations on the edges. As the set container is not thread-safe, to avoid race updates, we use OpenMP locks.

C. Updating the Data Structures

In our proposed fully dynamic algorithm, we maintain numerous integer values which are updated in parallel. We use the atomic built-ins from the GCC compiler for a race-free implementation. The *sup* of each tree edge is frequently updated by insert and delete operations. We use STL container *Sets* for faster updates. Additionally, we use OpenMP locks to avoid race conditions.

VI. EXPERIMENTAL RESULTS

In this section, we discuss the datasets and the configuration of the experimental platform. We further present the performance of our Incremental Dynamic algorithm over the static algorithms. Later, we discuss the performance of our Fully Dynamic Algorithm. Finally, we conclude by addressing the scalability of our algorithms.

A. Dataset and Experimental Platform:

In Table II, we provide the details of the datasets used in our experiments. We choose a healthy mix of datasets from different publicly available sources in road networks, web graphs, co-purchasing networks, and social networks. We use datasets with edge sizes varying from 1 M to 1 B edges in order to emphasize the scalability of our implementations.

For our experiments, we use a shared memory platform containing two 64 core AMD EPYC 7742 processors spread across two I/O hubs for a total of 128 cores. With simultaneous multi-threading, each core supports two threads of execution hence effectively providing 256 execution lanes. The EPYC processor is based on the Zen 2 microarchitecture from AMD and is built on the 7nm process. It has a base frequency of 2.24 GHz and a working frequency of 3.34 GHz. The processor has

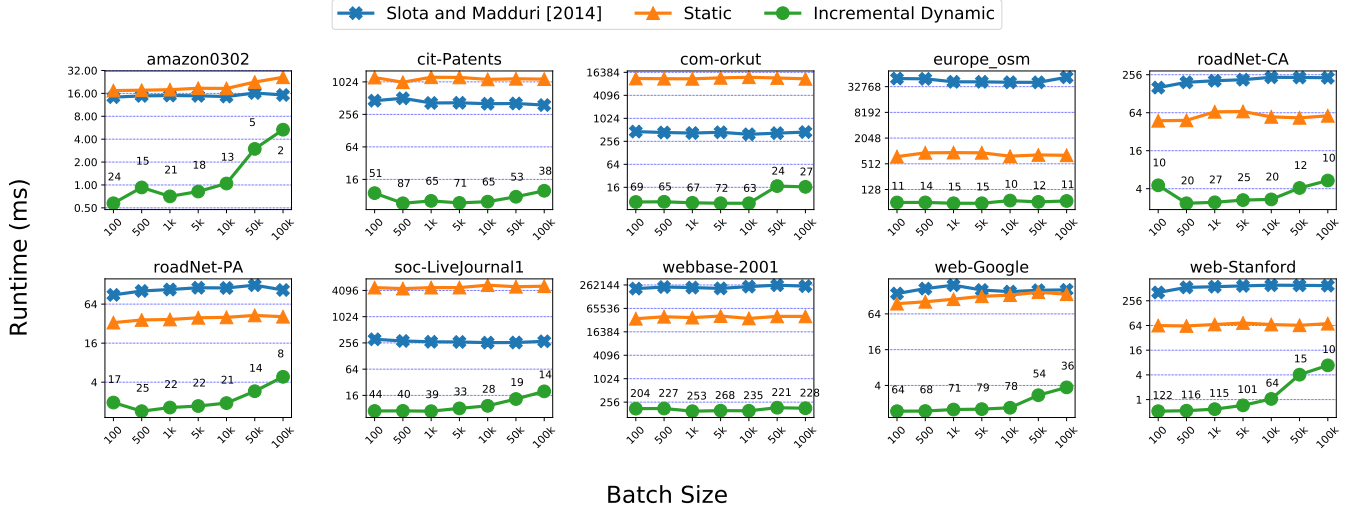


Fig. 3. Performance with varying batch size for insert operations shown in Algorithm 1. (Thread Count = 32)

a thermal design power (TDP) consumption of 225 watts and contains an L3 cache of 256 MB. The processor is connected to 1 TB DDR4 main memory spread across 8 NUMA nodes. The server is running CentOS 7. For the implementation, we compile our code with gcc version 7.5 with OpenMP version 4.5. We compile all programs with the $-O3$ flag.

We perform the comparative analysis of the multicore implementations of the proposed dynamic algorithms against the multicore implementations of the static algorithm by Slota *et al.* [17]², and our static algorithm (*Static*) described in Section III-B.

B. Performance of the Incremental Dynamic Algorithm (Algorithm 1)

We now study the overall performance of Algorithm 1 compared to static algorithms. For each of the graphs listed in Table II, we consider inserting edges in batches of size varying from 100 edges to 100,000 edges. From Figure 3, we can observe that the time taken by our Algorithm 1 and the static implementations of Slota *et al.* [17] and our static algorithm. The numbers on top of the line show the speedup of Algorithm 1 over the best of both static algorithms.

We notice that our incremental approach performs $58\times$ better on an average in comparison to the static implementations of Slota *et al.* [17], and our static algorithm. We attribute this speedup to the fact that our incremental algorithm processes only affected fundamental cycles in parallel. Moreover, the lock-free implementation increases the speedup.

C. Performance of the Fully Dynamic Algorithm

We study the performance of Algorithm 2 and 3 separately. From Section IV, we understand that deleting a tree edge is more expensive than deleting a non-tree edge. Precisely,

deletion of a tree edge involves the deletion and insertion of multiple non-tree edges. Thus, we segregate the queries into three kinds: Insert edges, Delete Non-tree edges, and Delete Tree edges. This partition of queries allows us to understand the performance gain in finer detail. The batch size varies from 50 to 1000 for non-tree edges and 25 to 200 for tree edges.

Insert Operations: Figure 4 shows the performance of the fully dynamic insert operations. We can see that our insert operation is on an average 8.63x better than the work by Slota *et al.* [17] and our static implementation.

Delete Non-Tree Edge Operations: In Figure 5 we show the performance of Algorithm 3 when a batch of non-tree edges are deleted. We observe that the average speedup obtained is 6.6x over the best of Slota *et al.* [17] and our static algorithm.

Delete Tree Edge Operations: Figure 6, shows the performance of the delete operations on tree edges specifically. We observe that in this case our average speedup is 2.55x. As tree-edge deletion is an expensive operation, large batch sizes are not practical.

From our experiments, we note that our implementations of our fully dynamic algorithms support a throughput of 63K edge insertions per second, 42K delete non-tree edges per second, and 1K delete tree edges per second.

D. Scalability Analysis

In Figure 7, we can observe the scaling results for the incremental technique shown in Algorithm III. We can see the performance of our method scales linearly up to 64 threads, where the performance improves from 705.4 ms (for 4 threads) to 134.01 ms (at 64 threads) for the largest graph (webbase) used in our experiments when the batch size is 10K. On an average the performance of all the graphs improves by 55.57% when run using 64 threads in comparison to 4 threads. The performance however drops after 64 threads due to the

²Code from Authors - <https://www.cs.rpi.edu/~slotag/publications.html>

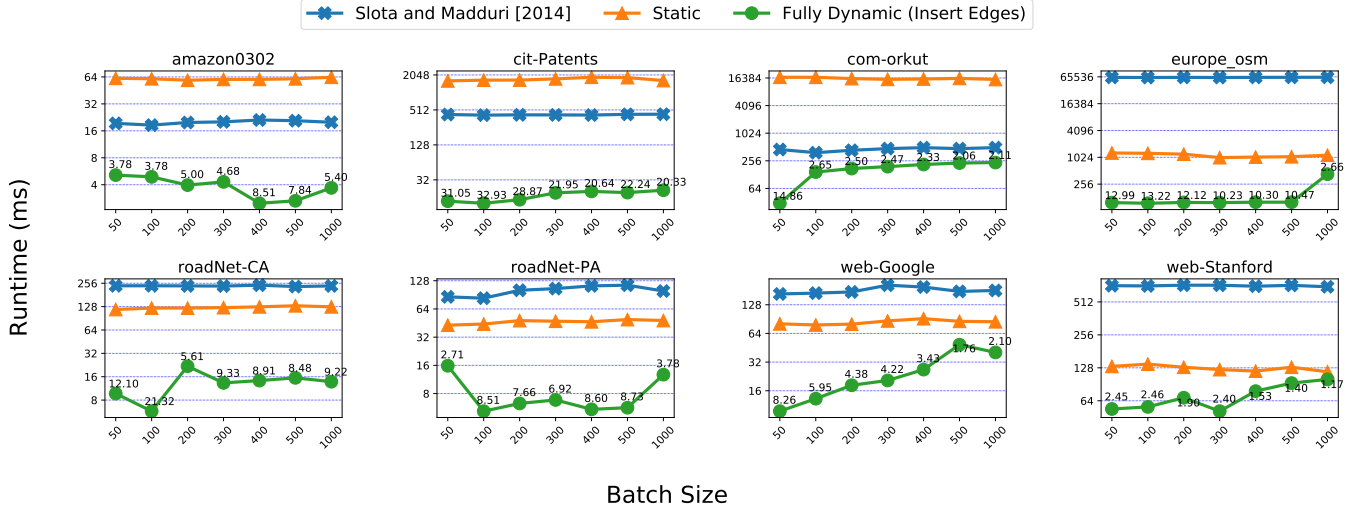


Fig. 4. Performance with varying batch size for insert operations shown in Algorithm 2. (Thread Count = 32)

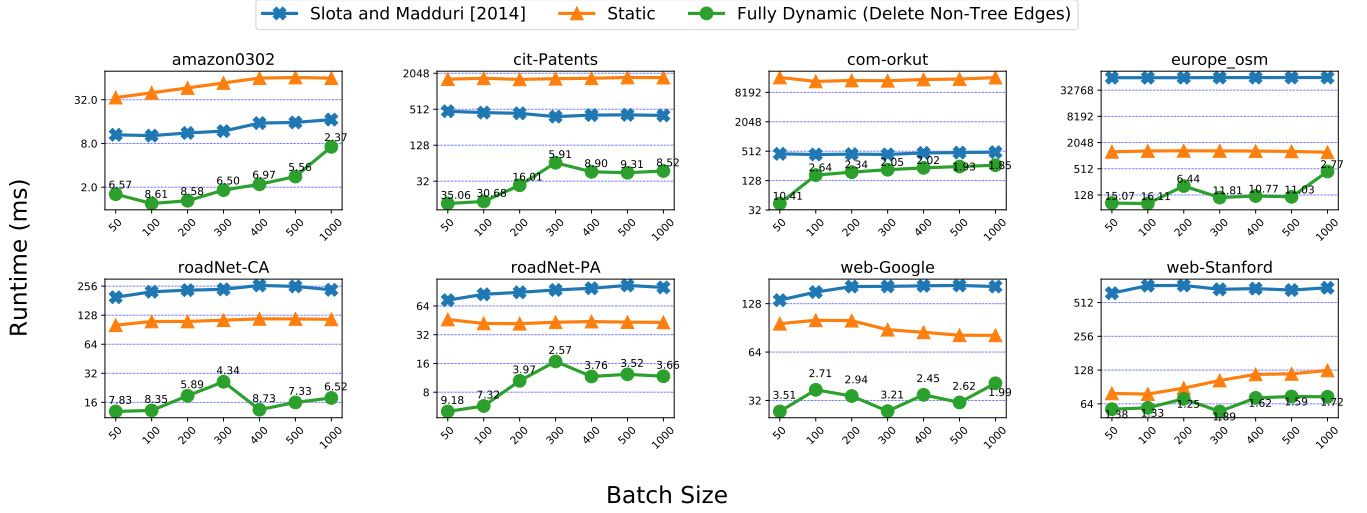


Fig. 5. Performance with varying batch size for deleting non-tree edges shown in Algorithm 3. (Thread Count = 32)

NUMA effect since we work on a platform where 128 compute cores are spread across two NUMA hubs. The improvement in scaling can be attributed to the parallel connected components computation, parallel processing of affected fundamental cycles, and the identification of the cut-edges and cut-vertices in parallel. Since the computations are fairly data parallel, a linear gain in performance is obtained when increasing the number of threads. We observe a similar trend for the fully dynamic updates also.

VII. CONCLUSION

We have designed parallel algorithms to update the cut vertices, cut edges and the biconnected components of a graph,

when a batch of edges are inserted or deleted dynamically. We have examined the running times of our incremental dynamic algorithm and fully dynamic algorithms against the best of Slota *et al.* implementation and our own static implementation, and the average speedup ranges from $12\times$ to $233\times$, and $1.4\times$ to $16.78\times$. The primary reason for this improvement is that, our algorithms traverse only affected fundamental cycles, and update the necessary LCA-graphs to finally update the biconnected components.

REFERENCES

- [1] T. A. Davis and Y. Hu, “The university of florida sparse matrix collection,” *ACM Trans. Math. Softw.*, pp. 1–25, Dec. 2011.
- [2] D. Eppstein, Z. Galil, G. F. Italiano, and A. Nissenzeig,

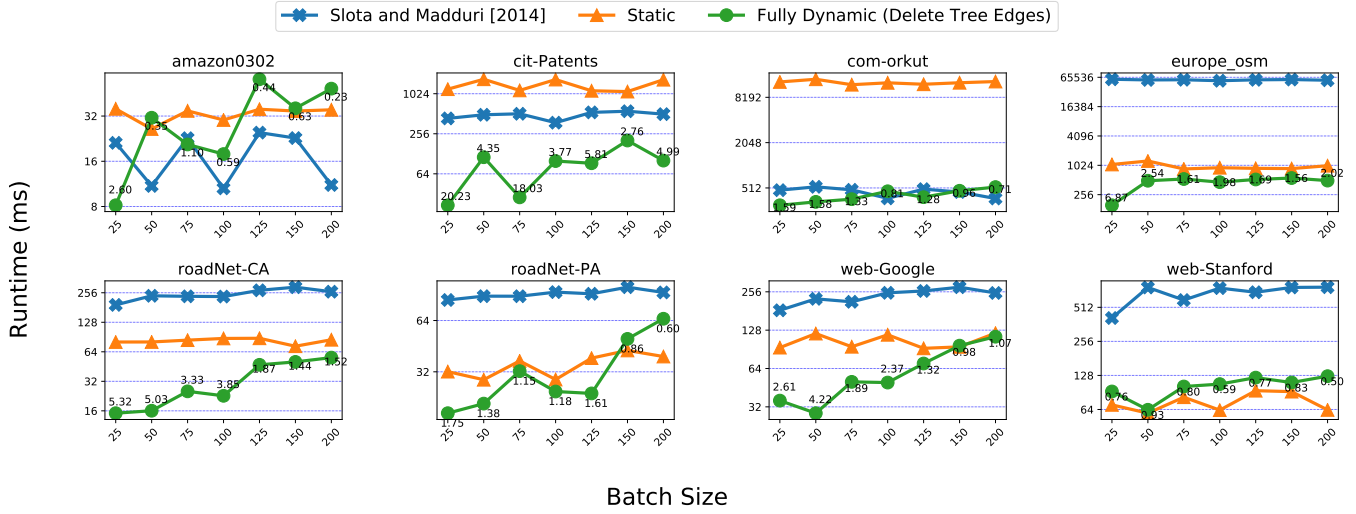


Fig. 6. Performance with varying batch size for deleting tree edges shown in Algorithm 3 (Thread Count = 32)

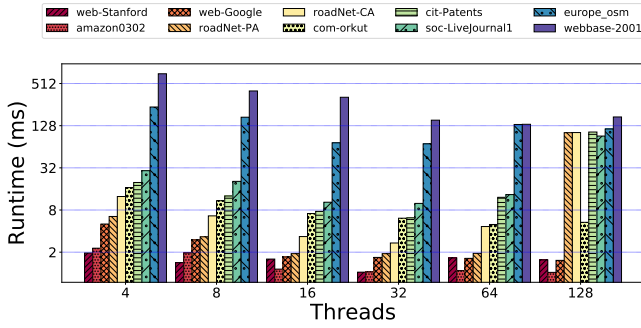


Fig. 7. Variation with different number of threads for the incremental updates.

“Sparsification — A technique for speeding up dynamic graph algorithms,” *J. ACM*, vol. 44, no. 5, pp. 669–696, September 1997.

- [3] Z. Galil and G. F. Italiano, “Fully dynamic algorithms for 2-edge connectivity,” *SIAM J. Comput.*, vol. 21, no. 6, pp. 1047–1069, 1992.
- [4] O. Green, R. McColl, and D. A. Bader, “A fast algorithm for streaming betweenness centrality,” in *2012 IEEE SocialCom 2012*, 2012, pp. 11–20.
- [5] M. Halappanavar, H. Lu, A. Kalyanaraman, and A. Tumeo, “Scalable static and dynamic community detection using Grapolo,” in *IEEE HPEC*, 2017, pp. 1–6.
- [6] M. R. Henzinger and V. King, “Randomized fully dynamic graph algorithms with polylogarithmic time per operation,” *J. ACM*, vol. 46, no. 4, pp. 502–516, 1999.
- [7] D. S. Hochbaum, “Why should biconnected components be identified first,” *Discrete Applied Mathematics*, vol. 42, pp. 203–210, 1993.
- [8] F. T. Jamour, S. Skiadopoulos, and P. Kalnis, “Parallel algorithm for incremental betweenness centrality on large graphs,” *IEEE TPDS*, vol. 29, no. 3, pp. 659–672, 2018.
- [9] A. Khanda, S. Srinivasan, S. Bhowmick, B. Norris, and S. K. Das, “A Parallel Algorithm Template for Updating Single-

Source Shortest Paths in Large-Scale Dynamic Networks,” *IEEE TPDS*, 2021.

- [10] J. Leskovec and A. Krevl, “SNAP Datasets: Stanford large network dataset collection,” <http://snap.stanford.edu/data>, Jun. 2014.
- [11] R. McColl, O. Green, and D. A. Bader, “A new parallel algorithm for connected components in dynamic graphs,” in *20th IEEE HiPC*, 2013, pp. 246–255.
- [12] —, “A new parallel algorithm for connected components in dynamic graphs,” in *20th IEEE HiPC*, 2013, pp. 246–255.
- [13] G. Ramalingam, *Bounded Incremental Computation*. Lecture Notes in Computer Science, Springer, 1996, vol. 1089.
- [14] A. E. Sariyüce, K. Kaya, E. Saule, and Ü. V. Çatalyürek, “Betweenness centrality on GPUs and heterogeneous architectures,” in *6th GPGPU Workshop*, 2013, pp. 76–85.
- [15] K. Shukla, S. C. Regunta, S. H. Tondomker, and K. Kothapalli, “Efficient parallel algorithms for betweenness- and closeness-centrality in dynamic graphs,” in *ACM ICS*, 2020, pp. 10:1–10:12.
- [16] N. Simsiri, K. Tangwongsan, S. Tirthapura, and K. Wu, “Work-efficient parallel union-find with applications to incremental graph connectivity,” in *Proc. Euro-Par*, 2016, pp. 561–573.
- [17] G. M. Slota and M. Kamesh, “Simple parallel biconnectivity algorithms for multicore platforms,” in *2014 21st HiPC*, 2014, pp. 1–10.
- [18] J. Soman, K. Kothapalli, and P. J. Narayanan, “A fast GPU algorithm for graph connectivity,” in *24th IEEE IPDPS Workshops*, 2010, pp. 1–8.
- [19] S. Srinivasan, S. Bhowmick, and S. K. Das, “Application of Graph Sparsification in Developing Parallel Algorithms for Updating Connected Components,” in *IPDPS Workshops 2016*. IEEE Computer Society, 2016, pp. 885–891.
- [20] S. Srinivasan, S. Pollard, S. Das, B. Norris, and S. Bhowmick, “A Shared-Memory Algorithm for Updating Tree-Based Properties of Large Dynamic Networks,” *IEEE Transactions on Big Data*, vol. PP, pp. 1–1, 09 2018.
- [21] Weifa Liang, R. P. Brent, and Hong Shen, “Fully dynamic maintenance of k-connectivity in parallel,” *IEEE TPDS*, vol. 12, no. 8, pp. 846–864, Aug 2001.